

# 15-441/641: Computer Networks

## Domain Name System

15-441 Spring 2019  
 Profs **Peter Steenkiste** & Justine Sherry




Fall 2019  
<https://computer-networks.github.io/sp19/>

**Carnegie  
 Mellon  
 University**

## Too Much of a Good Thing?

- Hosts have a
  - host name
  - IP address
  - MAC address
- There is a reason ..
- Remember?
- But how do we translate?



2

## IP to MAC Address Translation

- How does one find the Ethernet address of a IP host?
- Address Resolution Protocol - ARP
  - Broadcast search for IP address
    - E.g., "who-has 128.2.184.45 tell 128.2.206.138" sent to Ethernet broadcast (all FF address)
  - Destination responds (only to requester using unicast) with appropriate 48-bit Ethernet address
    - E.g., "reply 128.2.184.45 is-at 0:d0:bc:f2:18:58" sent to 0:c0:4f:d:ed:c6



3

## Caching ARP Entries

- Efficiency Concern
  - Would be very inefficient to use ARP request/reply every time need to send IP message to machine
- Each Host Maintains Cache of ARP Entries
  - Add entry to cache whenever you get ARP response
  - "Soft state": set timeout of ~20 minutes



4

## ARP Cache Example

- Show using command "arp -a"

```


Interface: 128.2.222.198 on Interface 0x1000003
Internet Address      Physical Address      Type
128.2.20.218          00-b0-8e-83-df-50     dynamic
128.2.102.129         00-b0-8e-83-df-50     dynamic
128.2.194.66          00-02-b3-8a-35-bf     dynamic
128.2.198.34          00-06-5b-f3-5f-42     dynamic
128.2.203.3           00-90-27-3c-41-11     dynamic
128.2.203.61          08-00-20-a6-ba-2b     dynamic
128.2.205.192         00-60-08-1e-9b-fd     dynamic
128.2.206.125         00-d0-b7-c5-b3-f3     dynamic
128.2.206.139         00-a0-c9-98-2c-46     dynamic
128.2.222.180         08-00-20-a6-ba-c3     dynamic
128.2.242.182         08-00-20-a7-19-73     dynamic
128.2.254.36          00-b0-8e-83-df-50     dynamic
  
```



5


## Challenge: Broadcast!

- Overhead scales (roughly) as  $N^2$  for an  $N$  host network
  - $N$  host does an ARP broadcast for each (new) destination
  - Each broadcast is delivered to  $N$  hosts
- Remember the solution?
- Subnetting!
  - Break up network into networks connected by router
- Not always a good idea
  - Extra complexity, management overhead, cost, ...



## Proxy ARP

- Limit the scope of ARP requests/responses inside an L2
- Proxy ARP makes it look like one network:
  - Host1 in N1 sends ARP for host 2 in N2
  - Proxy ARP looks up MAC address
    - May require discovery using ARP
  - Responds to host 1's request
    - Acts as proxy
  - Also forwards packets to host1
    - Acts as a switch



## Host Names & Addresses

- Host addresses: *e.g.*, 169.229.131.109
  - a number used by protocols
  - conforms to network structure (the "where")
- Host names: *e.g.*, linux.andrew.cmu.edu
  - mnemonic name usable by humans
  - conforms to organizational structure (the "who")
- The Domain Name System (DNS) is how we map from one to the other
  - a **directory service** for hosts on the Internet



## Why bother?

- Convenience
  - Easier to remember [www.google.com](http://www.google.com) than 74.125.239.49
- Provides a level of indirection!
  - Decoupled names from addresses
  - Many uses beyond just naming a specific host



## DNS provides Indirection

- Addresses can **change** underneath
  - Move [www.cnn.com](http://www.cnn.com) to a new IP address
  - Humans/apps are unaffected
- Name could map to **multiple** IP addresses
  - Enables load-balancing
- **Multiple names** for the same address
  - E.g., many services (mail, www, ftp) on same machine
- Allowing “host” names to evolve into “service” names



## DNS: Early days

- Mappings stored in a hosts.txt file (in /etc/hosts)
  - maintained by the Stanford Research Institute (SRI)
  - new versions periodically copied from SRI (via FTP)
- As the Internet grew this system broke down
  - SRI couldn't handle the load
  - conflicts in selecting names
  - hosts had inaccurate copies of hosts.txt
- The Domain Name System (DNS) was invented to fix this



## Obvious Solutions (1)

### Why not centralize DNS?

- Distant centralized database
  - Traffic volume
- Single point of failure
- Single point of update
- Single point of control
- Doesn't *scale*!



## Goals?

- Scalable
  - many names
  - many updates
  - many users creating names
  - many users looking up names
- Highly available
- Correct
  - no naming conflicts (uniqueness)
  - consistency
- Lookups are fast



## How?

- Partition the namespace – Hierarchy!
  - Distribute administration of each partition
    - Autonomy to update my own (machines') names
      - Translation of cmu.edu names is done by CMU
    - Don't have to track everybody's updates
  - Distribute name resolution for each partition
- *How should we partition things?*




## Key idea: hierarchical distribution

Three intertwined hierarchies

- Hierarchical namespace
  - As opposed to original flat namespace
- Hierarchically administered
  - As opposed to centralized administrator
- Hierarchy of servers
  - As opposed to centralized storage



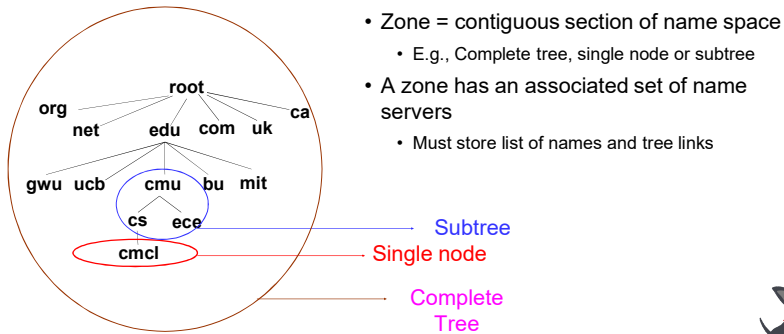
## DNS Design: Hierarchy Definitions



- Each node in hierarchy stores a list of names that end with same suffix
  - Suffix = path up tree
- E.g., given this tree, where would following be stored:
  - Fred.com
  - Fred.edu
  - Fred.cmu.edu
  - Fred.cmcl.cs.cmu.edu
  - Fred.cs.mit.edu



## DNS Design: Zone Definitions



17

## Server Hierarchy

- Top of hierarchy: Root servers
  - Location hardwired into other servers
- Next Level: Top-level domain (TLD) servers
  - .com, .edu, .uk, etc.
  - Managed professionally
- Bottom Level: **Authoritative** DNS servers
  - Actually store the name-to-address mapping
  - Maintained by the corresponding administrative authority

New TLDs started in 2012  
... expect to see more  
in the future.



## Server Hierarchy

- Every server knows the address of the root name server
  - Root servers know the address of all TLD servers
  - ...
  - An authoritative DNS server stores name-to-address mappings ("resource records") for all DNS names in the domain that it has authority for
- Each server stores a subset of the total DNS database
- Each server can discover the server(s) responsible for any portion of the hierarchy



## DNS Root

- Located in Virginia, USA



## DNS Root Servers

- 13 root servers (labeled A-M; see <http://www.root-servers.org/>)



## DNS Root Servers

- 13 root servers (labeled A-M; see <http://www.root-servers.org/>)
- Replicated via **any-casting**



## Anycast in a nutshell

- Routing finds shortest paths to destination
- What happens if multiple machines advertise the same address?
- The network will deliver the packet to the closest machine with that address
- This is called "anycast"
  - Very robust
  - Requires no modification to routing algorithms



## Programmer's View of DNS

- Conceptually, programmers can view the DNS database as a collection of millions of *host entry structures*:

```
/* DNS host entry structure */
struct addrinfo {
    int    ai_family; /* host address type (AF_INET) */
    size_t ai_addrlen; /* length of an address, in bytes */
    struct sockaddr *ai_addr; /* address */
    char *ai_canonname; /* official domain name of host */
    struct addrinfo *ai_next; /* other entries for host */
};
```

- Functions for retrieving host entries from DNS:
  - `getaddrinfo`: query key is a DNS host name.
  - `getnameinfo`: query key is an IP address.



## Properties of DNS Host Entries

- Different kinds of mappings are possible:
  - Simple case: 1-1 mapping between domain name and IP addr:
    - kittyhawk.cmcl.cs.cmu.edu maps to 128.2.194.242
  - Multiple domain names maps to the same IP address:
    - eeecs.mit.edu and cs.mit.edu both map to 18.62.1.6
  - Single domain name maps to multiple IP addresses:
    - www.google.com maps to multiple IP addrs.
- Some valid domain names don't map to any IP address:
  - for example: cmcl.cs.cmu.edu



25

## DNS Records

RR format: (class, name, value, type, ttl)

- DB contains tuples called resource records (RRs)
  - Classes = Internet (IN), Chaosnet (CH), etc.
  - Each class defines value associated with type

### FOR IN class:

- |   |  |
|---|--|
| <ul style="list-style-type: none"> <li>• Type=A               <ul style="list-style-type: none"> <li>• <b>name</b> is hostname</li> <li>• <b>value</b> is IP address</li> </ul> </li> <li>• Type=NS               <ul style="list-style-type: none"> <li>• <b>name</b> is domain (e.g. foo.com)</li> <li>• <b>value</b> is name of authoritative name server for this domain</li> </ul> </li> </ul> | <ul style="list-style-type: none"> <li>• Type=CNAME               <ul style="list-style-type: none"> <li>• <b>name</b> is an alias name for some "canonical" (the real) name</li> <li>• <b>value</b> is canonical name</li> </ul> </li> <li>• Type=MX               <ul style="list-style-type: none"> <li>• <b>value</b> is hostname of mailserver associated with <b>name</b></li> </ul> </li> </ul> |
|---|--|



26

## Inserting RRs into DNS

- Example: you just created company "FooBar"
- You get a block of IP addresses from your ISP
  - say 212.44.9.128/25
- Register [foobar.com](http://foobar.com) at registrar (e.g., NameCheap)
  - Provide registrar with names and IP addresses of your authoritative name server(s)
  - Registrar inserts RR pairs into the **.com TLD** server:
    - ([foobar.com](http://foobar.com), [dns1.foobar.com](http://dns1.foobar.com), **NS**)
    - ([dns1.foobar.com](http://dns1.foobar.com), **212.44.9.129**, **A**)
- Store resource records in your server [dns1.foobar.com](http://dns1.foobar.com)
  - e.g., type A record for [www.foobar.com](http://www.foobar.com)
  - e.g., type MX record for [foobar.com](http://foobar.com)



## Using DNS (Client/App View)

- Two components
  - Local DNS servers
  - Resolver software on hosts
- Each host has a resolver
  - Typically a library that applications can link to
- Client application
  - Obtain DNS name (e.g., from URL)
  - Triggers DNS request to its local DNS server




## Servers/Resolvers


- Name servers: generally responsible for some zone
  - Answer queries about their zone
- Local DNS server ("default name server")
  - Answer queries about the local zone
  - Also do lookup of distant host names for local hosts
    - Can cache the response for other local hosts!
  - Clients configured with the default server's address or learn it via a host configuration protocol




29



30

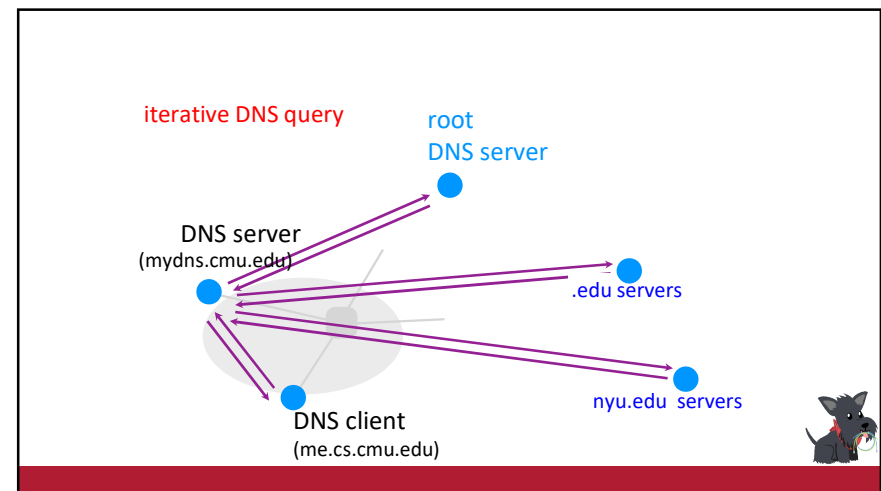
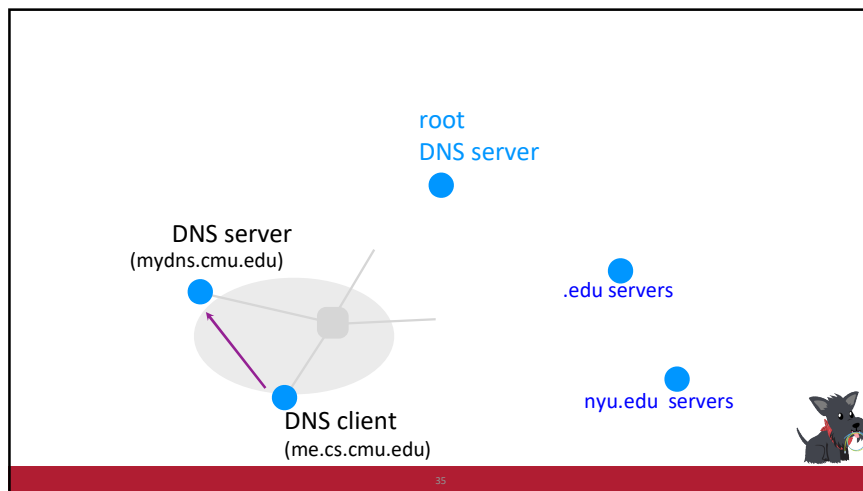
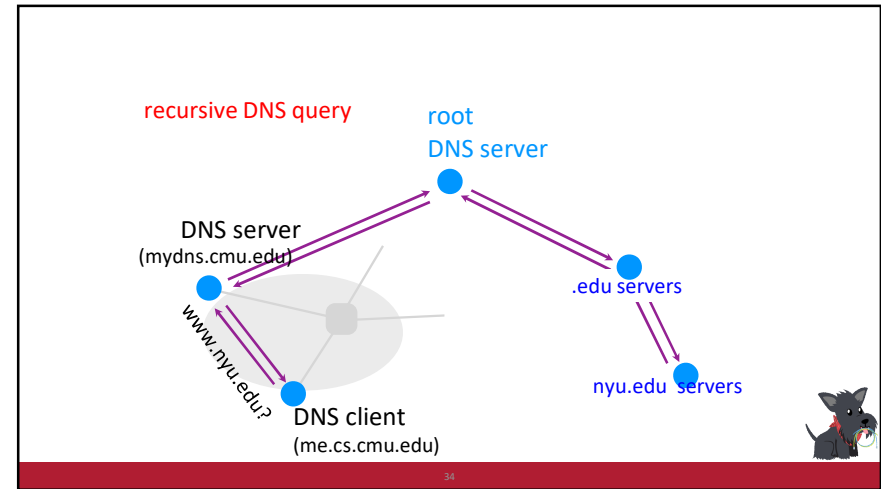
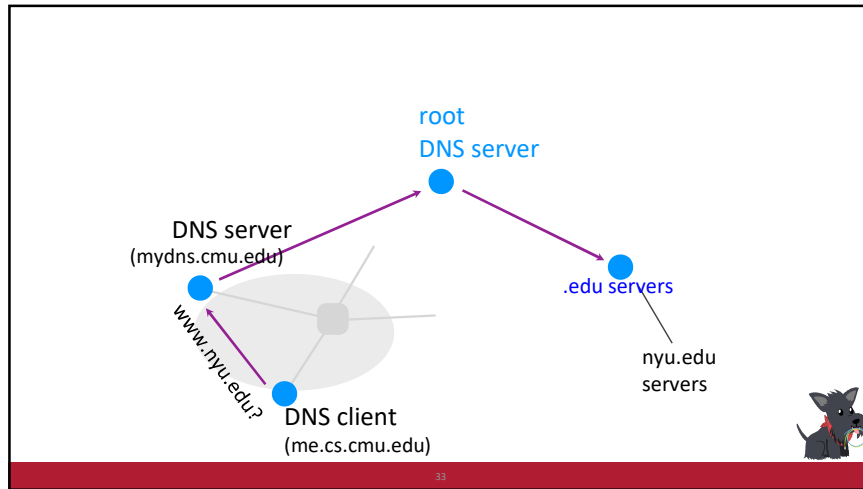


31



32





## Goals – how are we doing?

- Scalable
  - many names
  - many updates
  - many users creating names
  - many users looking up names
- Highly available



## Per-domain availability

- DNS servers are **replicated**
  - Primary and secondary name servers required
  - Name service available if at least one replica is up
  - Queries can be load-balanced between replicas
- Try alternate servers on timeout
  - **Exponential backoff** when retrying same server



## DNS Caching

- Caching of DNS responses at all levels
  - Reduces load at all levels
  - Reduces delay experienced by DNS client
- How DNS caching works
  - DNS servers cache responses to queries
  - Responses include a "time to live" (TTL) field
  - Server deletes cached entry after TTL expires
- Why caching is effective
  - The top-level servers very rarely change
  - Popular sites visited often → local DNS server often has the information cached



## Negative Caching

- Remember things that don't work
  - Misspellings like *www.cnn.comm* and *www.cnnn.com*
  - These can take a long time to fail the first time
  - Good to remember that they don't work
  - ... so the failure takes less time the next time around
- Negative caching is optional

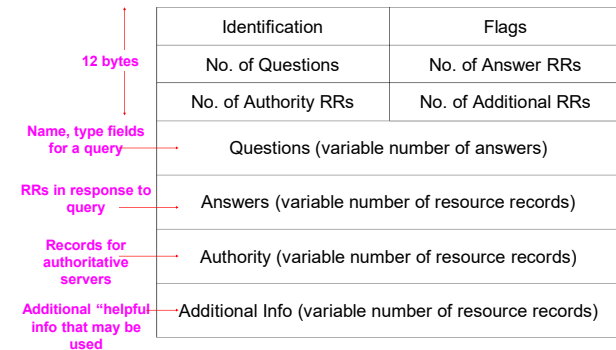


## Goals – how are we doing?

- Scalable
  - many names
  - many updates
  - many users creating names
  - many users looking up names
- Highly available
- Correct
  - no naming conflicts (uniqueness)
    - consistency
- Lookups are fast



## DNS Message Format



42



## DNS Header Fields

- Identification
  - Used to match up request/response
- Flags
  - 1-bit to mark query or response
  - 1-bit to mark authoritative or not
  - 1-bit to request recursive resolution
  - 1-bit to indicate support for recursive resolution




43

## How can one attack DNS?




44

- Impersonate the local DNS server
  - give the wrong IP address to the DNS client




45

- Impersonate the local DNS server
  - give the wrong IP address to the DNS client



46

- Poison the cache of a DNS server
  - trick the server into caching the wrong IP address



## How can one attack DNS?

- Impersonate the local DNS server
  - give the wrong IP address to the DNS client
- Denial-of-service the root or TLD servers
  - make them unavailable to the rest of the world
- Poison the cache of a DNS server
  - trick the server into caching the wrong IP address



## Enter: DNSSEC

Extension to DNS to improve DNS security.



49

## Enter DNSSEC

Extension to DNS to improve DNS security

- provides message authentication and integrity verification through cryptographic signatures
  - You know who provided the signature
  - No modifications between signing and validation
- It does not provide authorization
- It does not provide confidentiality
- It does not provide protection against DDOS




## DNSSEC: Deployment Status

- 89% of top-level domains (TLDs) zones signed.
- ~47% of country-code TLDs (ccTLDs) signed.
- Second-level domains (SLDs) vary widely:
  - Over 2.5 million .nl domains signed (~45%) (Netherlands). [\[1\]](#)
  - ~88% of measured zones in .gov are signed.
  - Over 50% of .cz (Czech Republic) domains signed.
  - ~24% of .br domains signed (Brazil). [\[2\]](#)
- While only about 0.5% of zones in .com are signed, that percentage represents ~600,000 zones.



51

## DNSSEC: Deployment Status



52

## Important Properties of DNS

- Easy unique, human-readable naming
- Hierarchy helps with scalability
- Caching lends scalability, performance
- Not strongly consistent
- Trust model has some problems!

