

The Internet End-End The Web

15-441 Spring 2018
Profs **Peter Steenkiste** & Justine Sherry



Thanks to Scott Shenker, Sylvia Ratnasamay, Peter Steenkiste, and Srinu Seshan for slides.

**Carnegie
Mellon
University**


1945: Vannevar Bush



- “As we may think”, Atlantic Monthly, July, 1945.
- Describes the idea of a distributed hypertext system
- A “memex” that mimics the “web of trails” in our minds



Dec 9, 1968: “The Mother of All Demos”



First demonstration of Memex-inspired system

Working prototype with hypertext, linking, use of a mouse...

<https://www.youtube.com/watch?v=74c8LntW7fo>



Many other iterations before we got to the World Wide Web

- MINITEL in France. <https://en.wikipedia.org/wiki/Minitel>
- Project Xanadu. https://en.wikipedia.org/wiki/Project_Xanadu
- (Note that you don't need to know any of this history for exams, this is just for the curious...)



1989: Tim Berners-Lee

1989: Tim Berners-Lee (CERN) writes internal proposal to develop a distributed hypertext system


- Connects “a web of notes with links”.
- Intended to help CERN physicists in large projects share and manage information

1990: TBL writes graphical browser for Next machines

1992-1994: NCSA/Mosaic/Netscape browser release



Lots of Traffic!



What is an Exabyte?

	10 ^x	10 ^x
Kilo	3	10
Mega	6	20
Giga	9	30
Tera	12	40
Peta	15	50
Exa	18	60
Zetta	21	70
Yotta	24	80

Network 1,000,000,000,000,000 Bytes
Storage 1,099,511,627,776 MByte

— A few years ago
— Today
— In a few years




Hyper Text Transfer Protocol (HTTP)

- Client-server architecture
 - Server is “always on” and “well known”
 - Clients initiate contact to server
- Synchronous request/reply protocol
 - Runs over TCP, Port 80
- Stateless
- ASCII format



Steps in HTTP Request/Response



Client-to-Server Communication

- HTTP Request Message

- Request line: method, resource, and protocol version
- Request headers: provide information or modify request
- Body: optional data (e.g., to "POST" data to the server)

```

request line  GET /somedir/page.html HTTP/1.1
header lines  Host: www.someschool.edu
              User-agent: Mozilla/4.0
              Connection: close
              Accept-language: fr
              (blank line)
carriage return line feed indicates end of message
  
```

The diagram shows an example of an HTTP request message. The 'request line' is 'GET /somedir/page.html HTTP/1.1'. The 'header lines' include 'Host: www.someschool.edu', 'User-agent: Mozilla/4.0', 'Connection: close', and 'Accept-language: fr'. A '(blank line)' follows the headers. A note indicates that the 'carriage return line feed' indicates the end of the message.



Server-to-Client Communication

- HTTP Response Message

- Status line: protocol version, status code, status phrase
- Response headers: provide information
- Body: optional data

```

status line  HTTP/1.1 200 OK
             (protocol, status code, status phrase)
header lines  Connection: close
             Date: Thu, 06 Aug 2006 12:00:15 GMT
             Server: Apache/1.3.0 (Unix)
             Last-Modified: Mon, 22 Jun 2006 ...
             Content-Length: 6821
             Content-Type: text/html
             (blank line)
data         data data data data data ...
             e.g., requested HTML file
  
```

The diagram shows an example of an HTTP response message. The 'status line' is 'HTTP/1.1 200 OK'. The 'header lines' include 'Connection: close', 'Date: Thu, 06 Aug 2006 12:00:15 GMT', 'Server: Apache/1.3.0 (Unix)', 'Last-Modified: Mon, 22 Jun 2006 ...', 'Content-Length: 6821', and 'Content-Type: text/html'. A '(blank line)' follows the headers. The 'data' section contains 'data data data data data ...', which is noted as 'e.g., requested HTML file'.




HTTP is *Stateless*

- Each request-response treated independently
 - Servers *not* required to retain state
- **Good:** Improves scalability on the server-side
 - Failure handling is easier
 - Can handle higher rate of requests
 - Order of requests doesn't matter
- **Bad:** Some applications **need** persistent state
 - Need to uniquely identify user or store temporary info
 - e.g., Shopping cart, user profiles, usage tracking, ...



How to Maintain State in a Stateless Protocol: Cookies

- *Client-side* state maintenance
 - Client stores small state on behalf of server
 - Client sends state in future requests to the server
- Can provide authentication



Performance Issues

Performance Goals

- User
 - fast downloads (not identical to low-latency commn..!)
 - high availability
- Content provider
 - happy users (hence, above)
 - cost-effective infrastructure
- Network (secondary)
 - avoid overload

Solutions?

- User
 - fast downloads (not identical to low-latency commn..!)
 - high availability
- Content provider
 - happy users (hence, above)
 - cost-effective delivery infrastructure
- Network (secondary)
 - avoid overload

Improve HTTP to
compensate for
TCP's weak spots

Solutions?

- User
 - fast downloads (not identical to low-latency commn.!!)
 - high availability
- Content provider
 - happy users (hence, above)
 - cost-effective delivery infrastructure
- Network (secondary)
 - avoid overload

Improve HTTP to
compensate for
TCP's weak spots

Caching and Replication



Solutions?

- User
 - fast downloads (not identical to low-latency commn.!!)
 - high availability
- Content provider
 - happy users (hence, above)
 - cost-effective delivery infrastructure
- Network (secondary)
 - avoid overload

Improve HTTP to
compensate for
TCP's weak spots

Caching and Replication

Exploit economies of scale
(Webhosting, CDNs, datacenters)



HTTP Performance

- Most Web pages have multiple objects
 - e.g., HTML file and a bunch of embedded images
- How do you retrieve those objects (naively)?
 - One item at a time, i.e., one "GET" per TCP connection
 - Solution used in HTTP 0.9, and 1
- New TCP connection per (small) object!
 - Lots of handshakes
 - Congestion control state lost across connections



Typical Workload (Web Pages)

- Multiple (typically small) objects per page
- File sizes
 - Heavy-tailed
 - Pareto distribution for tail
 - Lognormal for body of distribution
- Embedded references
 - Number of embedded objects also Pareto
- This plays havoc with performance. Why?
- Solutions?

• Lots of small objects versus TCP
• 3-way handshake
• Lots of slow starts
• Extra connection state


Improving HTTP Performance: Persistent Connections

- Maintain TCP connection across multiple requests
 - Including transfers subsequent to current page
 - Client or server can tear down connection
- Performance advantages:
 - Avoid overhead of connection set-up and tear-down
 - Allow TCP to learn more accurate RTT estimate
 - Allow TCP congestion window to increase
 - i.e., leverage previously discovered bandwidth
- Drawback? Head of line blocking
 - A "slow object" blocks retrieval of all later requests, including "fast" objects
- Default in HTTP/1.1



Improving HTTP Performance: Pipelined Requests & Responses


- Batch requests and responses to reduce the number of packets
- Multiple requests can be contained in one TCP segment
- Head of line blocking issues remains: a delay in Transfer 2 delays all later transfers



Improving HTTP Performance: Concurrent Requests & Responses

- Use multiple connections *in parallel*
- Speeds up retrieval by $\sim m$
- Does not necessarily maintain order of responses
- Partially deals with HOL blocking

- Client =
- Content provider =
- Network = Why?



Scorecard: Getting n Small Objects

Time dominated by latency

- One-at-a-time: $\sim 2n$ RTT
- M concurrent: $\sim 2[n/m]$ RTT
- Persistent: $\sim (n+1)$ RTT
- Pipelined: ~ 2 RTT
- Pipelined/Persistent: ~ 2 RTT first time, RTT later



Scorecard: Getting n Large Objects

Time dominated by bandwidth

- One-at-a-time: $\sim nF/B$
- M concurrent: $\sim [n/m] F/B$
 - assuming shared with large population of users
 - and each TCP connection gets the same bandwidth
- Pipelined and/or persistent: $\sim nF/B$
 - The only thing that helps is getting more bandwidth..




Improving HTTP Performance: Caching

- Why does caching work?
 - Exploits *locality of reference*
- How well does caching work?
 - Very well, up to a limit
 - Large overlap in content
 - But many unique requests
- Trend: increase in dynamic content
 - E.g., customizing of web pages
 - Reduces benefits of caching
 - Some exceptions, e.g., video




Improving HTTP Performance: Caching: Where?

- Baseline: Many clients transfer same information
 - Generate unnecessary server and network load
 - Clients experience unnecessary latency
- Everywhere!
 - Client
 - Forward proxies
 - Reverse proxies
 - Content Distribution Network




Improving HTTP Performance: Caching: Clients

- Clients keep a local cache of recently accessed objects
 - Clients often have a small number of web pages they access frequently
 - Leads to reuse of logos, old content, java scripts, ...
- Cheap: no additional infrastructure needed
- But caching closer to server can lead to higher hit rates!




Improving HTTP Performance: Caching with Reverse Proxies

- Cache documents close to **server**
→ decrease server load
- Typically done by content provider



Improving HTTP Performance: Caching with Forward Proxies

- Cache documents close to **clients**
→ decrease latency
- Typically done by ISPs or enterprises
→ reduce provider traffic load



Improving HTTP Performance: Caching: How to Avoid Stale Content

- Modifier to GET requests:
 - **If-modified-since** – returns “not modified” if resource not modified since specified time

```
GET /~ee122/fa13/ HTTP/1.1
Host: inst.eecs.berkeley.edu
User-Agent: Mozilla/4.03
If-modified-since: Sun, 27 Oct 2013 22:25:50 GMT
<CRLF>
```

- Client specifies “if-modified-since” time in request
- Server compares this against “last modified” time of resource
- Server returns “Not Modified” if resource has not changed
- or a “OK” with the latest version otherwise

Improving HTTP Performance: Caching: Helping the Cache

- Modifier to GET requests:
 - **If-modified-since** – returns “not modified” if resource not modified since specified time
- Response header:
 - **Expires** – how long it's safe to cache the resource
 - **No-cache** – ignore all caches; always get resource directly from server

Improving HTTP Performance: Replication

- Replicate popular Web site across many machines
 - Spreads load on servers
 - Places content closer to clients
 - Helps when content isn't cacheable
- Problem: Want to direct client to particular replica
 - Balance load across server replicas
 - Pair clients with nearby servers
- Common solution:
 - DNS returns different addresses based on client's geo location, server load, etc.



Improving HTTP Performance: Content Distribution Networks

- Caching and replication as a service
- Large-scale distributed storage infrastructure (usually administered by one entity)
 - e.g., Akamai has servers in 20,000+ locations
- Combination of (pull) caching and (push) replication
 - **Pull:** Direct result of clients' requests
 - **Push:** Expectation of high access rate
- Also do some processing
 - Handle *dynamic* web pages
 - *Transcoding*



Recall: CDN Example – Akamai

- Akamai creates new domain names for each client
 - e.g., a128.g.akamai.net for cnn.com
- The CDN's DNS servers are authoritative for the new domains
- The client content provider modifies its content so that embedded URLs reference the new domains.
 - "Akamaize" content
 - e.g.: <http://www.cnn.com/image-of-the-day.gif> becomes <http://a128.g.akamai.net/image-of-the-day.gif>
- Requests for embedded objects are sent to CDN's infrastructure...



Cost-Effective Content Delivery

- General theme: multiple sites hosted on shared physical infrastructure
 - efficiency of statistical multiplexing
 - economies of scale (volume pricing, etc.)
 - amortization of human operator costs
- Examples:
 - Web hosting companies
 - CDNs
 - Cloud infrastructure



Performance Issues

Are We Done Yet?


Some Challenges with HTTP 1.1

- Head of line blocking: “slow” objects delay later requests
 - E.g., objects from remote storage versus objects in local memory
- Browsers open multiple TCP connections to achieve parallel transfers
 - Increases throughput and reduces impact HOL blocking
 - Increases load on servers and network
- HTTP headers are big
 - Cost higher for small objects
- Objects have dependencies, different priorities
 - Javascript versus images
 - Extra RTTs for “dependent” objects



40

Example of Head of Line Blocking



Source: <http://chimera.labs.oreilly.com/books/12300/0000345/ch11.html>


HTTP 2.0 to the Rescue

- Responses are multiplexed over single TCP connection
 - Server can send response data whenever it is ready
 - “Fast” objects can bypass slow objects – avoids HOL blocking
 - Fewer handshakes, more traffic (help cong. ctrl., e.g., drop tail)
- Multiplexing uses prioritized flow controlled streams
 - Urgent responses can bypass non-critical responses
 - ≈ multiple parallel prioritized TCP connections, but over one TCP connection
- HTTP headers are compressed
- A PUSH feature allows server to push embedded objects to the client without waiting for a client request
 - Avoids an RTT
- Default is to use TLS – fall back on 1.1 otherwise



42

HTTP/2 Multi-Streams Multiplexing



Bit	+0..7	+8..15	+16..23	+24..31
0	Length			Type
32	Flags			
40	R	Stream Identifier		
...	Frame Payload			

HTTP/2 Binary Framing

<https://tools.ietf.org/html/rfc7540>




Multiplexing

- Traffic sent as frames over prioritized streams
- Frames types: headers, data, settings, window updates and push promise
- Sender sends high priority frames first
- Frames are pulled from a per-stream queue when TCP is ready to accept more data
- Reduces queueing delay
- Each stream is flow controlled
- Receiver opens window faster for high priority streams
- Replicates TCP function but at finer granularity
- Clearly adds complexity to HTTP library



HTTP/2 Server Push



stream 1: /page.html (client request)
 stream 2: /script.js (push promise)
 stream 4: /style.css (push promise)



HTTP 2 PUSH Features

- Server can “push” objects that it knows (or thinks) the client will need
- Avoids delay of having client parse the page and requesting the objects (> RTT)
- But what happens if object is in the client cache – Oops!
- Server sends PUSH_PROMISE before the PUSH
- Client can cancel/abort the PUSH
- How does server know what to PUSH?
- Very difficult problem with dynamic content
- Javascripts can rewrite web page – changes URLs
- Also: benefits limited to objects from the origin server



Next Tuesday: Midterm Review

Use Piazza to request topics
Use midterm_review folder

